svmodt: An R Package for Linear SVM-Based Oblique Decision Trees

by Aneesh Agarwal, Jack Jewson, and Erik Sverdrup

Abstract An abstract of less than 150 words.

1 Introduction

2 Background

Decision Trees

Decision Trees (DTs) are interpretable classification models that represent their decision-making process through a hierarchical, tree-like structure. This structure comprises internal nodes containing splitting criteria and terminal (leaf) nodes corresponding to class labels. The nodes are connected by directed edges, each representing a possible outcome of a splitting criterion. Formally, a DT can be expressed as a rooted, directed tree $T = (G(V, E), v_1)$, where V denotes the set of nodes, E represents the set of edges linking these nodes, and v_1 is the root node.

If the tree T has m nodes, then for any $j \in \{1, \ldots, m\}$, the set of child nodes of $v_j \in V$ can be defined as:

$$N^+(v_j) = \{v_k \in V \mid k \in \{1, \ldots, m\}, \ k \neq j, \ (v_j, v_k) \in E\}.$$

Here, $N^+(v_j)$ denotes the set of nodes that are directly connected to v_j through outgoing edges, representing all possible child nodes that can be reached from v_j within the tree structure (Rivera-Lopez and Canul-Reich, 2018).

Decision tree algorithms can be categorized based on whether the same type of test is applied at all internal nodes. **Homogeneous trees** employ a single algorithm throughout (e.g., univariate or multivariate splits), whereas **hybrid trees** allow different algorithms such as linear discriminant functions, k-nearest neighbors, or univariate splits that can be used in different subtrees (Brodley and Utgoff, 1995). Hybrid trees exploit the principle of *selective superiority*, allowing subsets of the data to be modeled by the most appropriate classifier, thereby improving flexibility and accuracy.

Univariate Decision Trees

Univariate Decision Trees (UDTs) trees represent axis-parallel hyperplanes dividing the instance space into several disjoint regions. Axis-parallel decision trees, such as CART and C4.5, represent two of the most widely used algorithms for classification tasks. The **CART (Classification and Regression Trees)** algorithm employs a binary recursive partitioning procedure capable of handling both continuous and categorical variables as predictors or targets. It operates directly on raw data without requiring binning. The tree is expanded recursively until no further splits are possible, after which **cost-complexity pruning** is applied to remove branches that contribute least to predictive performance. This pruning process generates a sequence of nested subtrees, from which the optimal model is selected using independent test data or cross-validation, rather than internal training measures (Breiman et al., 1984).

In contrast, **C4.5**, an extension of the earlier **ID3** algorithm (Quinlan, 1986), utilizing information theory measures such as **information gain** and **gain ratio** to select the most informative attribute for each split (Quinlan, 2014). C4.5 also includes mechanisms to handle missing attribute values by weighting instances according to the proportion of known data and employs an **error-based pruning** method to reduce overfitting. Although these techniques are effective across diverse datasets, studies have shown that the choice of pruning strategy and stopping criteria can significantly affect model performance across different domains (Mingers, 1989; Schaffer, 1992).

While UDTs are highly interpretability, they are characterised by several representational limitations. Such trees often grow unnecessarily large, as they must approximate complex relationship between features through multiple axis-aligned partitions. This can result in the replication of subtrees and repeated testing of the same feature along different paths, both of which reduce efficiency and hinder generalization performance (Pagallo and Haussler, 1990).

Multivariate Decision Trees

Multivariate decision trees (MDTs) extends UDTs by allowing each internal node to perform splits based on linear or nonlinear combinations of multiple features. This flexibility enables the tree to form oblique decision boundaries that more accurately partition the instance space. For example, a single multivariate test such as $x + y < 8$ can replace multiple univariate splits needed to approximate the same boundary. The construction of MDTs introduces several design considerations, including how to represent multivariate tests, determine their coefficients, select features to include, handle symbolic and missing data, and prune to avoid overfitting (Brodley and Utgoff, 1995).

Various optimization algorithms—such as recursive least squares (Young, 1984), the pocket algorithm (Gallant, 1986), or thermal training (Frean, 1990)—may be used to estimate the weights. However, MDTs trade interpretability for representational power and often require additional mechanisms for **local feature selection**, such as *sequential forward selection* (SFS) or *sequential backward elimination* (SBE) (Kittler, 1986).

Empirical comparisons across multiple datasets demonstrate that multivariate trees generally achieve higher accuracy and smaller tree sizes than their univariate counterparts, though this comes at the cost of reduced interpretability. Moreover, MDTs retain key advantages of standard decision trees—such as sequential split criteria evaluation and transparent decision procedures—while offering improved modeling flexibility for complex datasets (Koziol and Wozniak, 2009; Friedl and Brodley, 1997; Liu and Setiono, 1998; Cañete et al., 2021).

Support Vector Machines (SVMs)

Support Vector Machines (SVMs) are powerful supervised learning models used for classification and regression tasks. They aim to determine an optimal separating hyperplane that maximizes the margin between different classes in the data. This margin-based approach enhances the generalization ability of the model, making SVMs robust and effective for many real-world problems (Cristianini, 2000).

Linear SVMs

A simplest **linear SVMs** construct a separating hyperplane in an n-dimensional space such that the margin between the classes is maximized. Given a training dataset $\{(\mathbf{x}_i, y_i)\}_{i=1}^{N}$, where $\mathbf{x}_i \in \mathbb{R}^n$ and $y_i \in \{-1, +1\}$, the decision function is defined as:

$$f(\mathbf{x}) = \text{sign}(\mathbf{w}^\top \mathbf{x} + b),$$

where $\mathbf{w}$ is the weight vector perpendicular to the hyperplane, and b is the bias term. The optimal hyperplane is the one that maximizes the distance between the closest points of each class (the **support vectors**) and the hyperplane itself (Cortes and Vapnik, 1995).

Mathematically, the optimization problem for a hard-margin SVM (i.e., assuming the data are linearly separable) can be formulated as:

$$\min_{\mathbf{w}, b} \ \frac{1}{2} \|\mathbf{w}\|^2$$

subject to:

$$y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1, \quad i = 1, 2, \ldots, N.$$

Here, $\|\mathbf{w}\|$ represents the norm of the weight vector and acts as a regularization term that controls the complexity of the model. The constraint ensures that all data points are correctly classified and lie outside the margin boundaries.

However, in most practical situations, perfect linear separability is not achievable. To address this, **soft-margin SVMs** introduce slack variables $\xi_i \geq 0$ to allow certain violations of the margin constraints, resulting in the following optimization problem:

$$\min_{\mathbf{w}, b, \xi} \ \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^{N} \xi_i$$

subject to:

$$y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 - \xi_i, \quad \xi_i \geq 0, \quad i = 1, 2, \ldots, N.$$

The parameter $C > 0$ controls the trade-off between maximizing the margin and minimizing the classification error on the training data. A large C penalizes mis-classifications heavily, leading to a

Table 1: Commonly used kernel functions and their parameters.

Kernel Type	Mathematical Definition	Key Parameters
Linear	$K(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i^\top \mathbf{x}_j$	None
Polynomial	$K(\mathbf{x}_i, \mathbf{x}_j) = (\mathbf{x}_i^\top \mathbf{x}_j + 1)^d$	Degree d
Gaussian	$K(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(-\frac{\|\mathbf{x}_i - \mathbf{x}_j\|^2}{2\sigma^2}\right)$	Bandwidth σ
RBF	$K(\mathbf{x}_i, \mathbf{x}_j) = \exp(-\gamma\|\mathbf{x}_i - \mathbf{x}_j\|^2)$	γ
Sigmoid	$K(\mathbf{x}_i, \mathbf{x}_j) = \tanh(\kappa \mathbf{x}_i^\top \mathbf{x}_j + \theta)$	κ, θ

narrower margin, whereas a smaller C allows more flexibility, potentially improving generalization in the presence of noise.

The solution to this constrained optimization problem is obtained using **Lagrange multipliers**, resulting in a dual formulation expressed as:

$$\max_{\alpha} \sum_{i=1}^{N} \alpha_i - \frac{1}{2} \sum_{i=1}^{N} \sum_{j=1}^{N} \alpha_i \alpha_j y_i y_j \langle \mathbf{x}_i, \mathbf{x}_j \rangle$$

subject to:

$$\sum_{i=1}^{N} \alpha_i y_i = 0, \quad 0 \leq \alpha_i \leq C.$$

The data points corresponding to non-zero α_i values are the **support vectors**, which define the decision boundary. The resulting decision function for a new observation $\mathbf{x}$ is given by:

$$f(\mathbf{x}) = \text{sign}\left(\sum_{i=1}^{N} \alpha_i y_i \langle \mathbf{x}_i, \mathbf{x} \rangle + b\right).$$

This formulation highlights one of the most important properties of SVMs — the decision boundary depends only on a subset of the training data (the support vectors), making SVMs both efficient and robust in representing the learned model (Cervantes et al., 2020).

Non-Linear SVMs

While linear classifiers provide useful insights, they are often inadequate for real-world datasets, where classes are not linearly separable. In such cases, SVMs can be extended to create **nonlinear decision boundaries** by mapping the input vectors into a higher-dimensional **feature space** using a nonlinear transformation $\phi : \mathbb{R}^n \to \mathcal{F}$. The linear transformation is then achieved in the transformed space using:

$$f(\mathbf{x}) = \text{sign}(\mathbf{w}^\top \phi(\mathbf{x}) + b).$$

However, directly computing $\phi(\mathbf{x})$ can be computationally expensive. To address this, SVMs employ the **kernel trick**, where the dot product in the feature space is replaced with a kernel function $K(\mathbf{x}_i, \mathbf{x}_j)$:

$$K(\mathbf{x}_i, \mathbf{x}_j) = \langle \phi(\mathbf{x}_i), \phi(\mathbf{x}_j) \rangle.$$

The resulting decision function becomes:

$$f(\mathbf{x}) = \text{sign}\left(\sum_{i=1}^{N} \alpha_i y_i K(\mathbf{x}_i, \mathbf{x}) + b\right),$$

where α_i are the Lagrange multipliers obtained during training.

For a function K to be a valid kernel, it must satisfy **Mercer's condition** (Vapnik, 2013) i.e., the kernel matrix must be symmetric and positive semi-definite. Table 1

Although SVMs exhibit strong theoretical foundations and robust generalization capabilities, they present several practical limitations. Model performance is highly dependent on the appropriate selection of hyperparameters such as the regularization term (C) and kernel parameters (e.g., γ), which govern the trade-off between margin maximization and misclassification tolerance (Nanda et al., 2018). Training an SVM requires solving a quadratic programming (QP) optimization problem involving an $n \times n$ kernel matrix, where n denotes the number of training samples, leading to quadratic growth in both computational time and memory usage (Dong et al., 2005). This makes SVMs computationally

expensive for large-scale datasets. Moreover, SVMs are inherently designed for binary classification, necessitating decomposition strategies such as One-vs-One and One-vs-All for multi-class problems (Hsu and Lin, 2002). Their performance also tends to degrade in imbalanced data settings, where the decision boundary becomes biased toward the majority class (Cervantes et al., 2020).

Hybrid Decision Trees with Support Vector Machines

The earliest formalization of Support Vector Machine–based Decision Trees (DTSVMs) was introduced by Bennett and Blue (1998) , who extended the principles of Statistical Learning Theory (Vapnik, 2013) and Structural Risk Minimization (SRM) to the construction of binary decision trees. Each node of a decision tree was treated as an SVM that partitions data along an optimal hyperplane. This approach allowed for multivariate SVM decisions at each node. TAKAHASHI and Abe (2002) extended DTSVMs to multiclass problems, using a recursive tree-based partitioning where the root separates the most separable class (or classes) from the rest. Subsequent nodes repeat this process until each leaf contains a single class, ensuring that all regions of the feature space are classified.

Subsequent studies have built upon this foundation to address scalability, optimization, and generalization issues. Optimal Decision Tree SVM (ODT-SVM) (Bala and Agrawal, 2011) introduced split-selection criteria based on Gini index, information gain, and scatter matrix separability to balance tree interpretability and margin-based precision.

Oblique and Rotation-Based Extensions

Recent research has explored oblique and rotation-based decision ensembles that generalize the DTSVM concept. Oblique Double Random Forests with MPSVM (MPDRaF) (Ganaie et al., 2022) introduce multivariate (oblique) splits at each node using Multi-Plane SVM (MPSVM) formulations (Mangasarian and Wild, 2006), enhancing the geometric flexibility of decision boundaries. Regularization variants—such as Tikhonov, axis-parallel, and null-space—address sample-size limitations at deeper nodes, ensuring better generalization.

Similarly, Rotation-based Double Random Forests (DRaF) (Ganaie et al., 2022) employ Principal Component Analysis (PCA) and Linear Discriminant Analysis (LDA) transformations at non-terminal nodes to create diverse subspaces, improving ensemble diversity and classification stability.

Modern Multi-class Integration

Recent works such as the Oblique Decision Tree Ensemble (ODTE) (Montañana et al., 2025) embed multiple SVM classifiers (e.g., one-vs-one or one-vs-rest) directly at each split, dynamically selecting the model that minimizes class impurity. This enables efficient n-ary classification within a single tree structure, mitigating the scalability and imbalance limitations of traditional binary SVM extensions.

Related Software Packages

Several software packages provide functionality related to SVM-based decision trees, but fully-featured implementations of SVM Oblique Decision Trees (SVMODT) remain limited, particularly in R. In R, very scarce SVMODT implementation is available. The **PPforest** package implements projection pursuit random forests, where trees are constructed by splitting on linear combinations of variables to maximize class separation, supporting multi-class problems and variable importance measures. In Python, **STree** (Montañana et al., 2021) provides an oblique decision tree classifier based on SVM nodes, where each node is built and split using scikit-learn's **SVC** models (Pedregosa et al., 2011). Java-based frameworks such as **Weka** provide classical decision trees (**J48**) and SVMs (**SMO**) that can be adapted for hierarchical SVM trees (Menkovski et al., 2008). Overall, while multiple languages provide the building blocks for SVM training, decision tree construction, and ensemble methods— there is currently no widely available, fully integrated SVMODT software in R, highlighting the need for custom implementations or adaptation of existing tools for research and applied purposes.

3 Methodology

To better understand the structure and mechanics of decision trees, we first implemented a standard axis-parallel decision tree in R from scratch before extending it into the oblique (SVM-based) variant.

Custom Univariate Decision Tree with Integrated Cost-Complexity Pruning

This implementation presents a decision tree algorithm with **integrated cost-complexity pruning**, where pruning decisions are made during tree construction rather than as a post-processing step. The algorithm supports multiple splitting criteria and performs real-time pruning based on cost-complexity analysis.

Algorithm Description

At each node v with samples $\mathcal{D}_v$ and labels $\mathcal{Y}_v$, the algorithm:

1. **Compute Node Statistics**: At each node of the decision tree, several key statistics are computed to guide the splitting process. The class probabilities are calculated as $p_c = \frac{n_c}{n}$, where n_c is the count of samples belonging to class c in the node, and $n = |D_v|$ is the total number of samples at the node. The predicted class is assigned as the class with the highest probability,

$$\hat{y} = \arg\max_c p_c.$$

2. **Select Optimal Split**: The optimal split at each node is determined using one of three criteria:

   - **Gini Impurity** is defined as

   $$\text{Gini}(\mathcal{Y}) = 1 - \sum_{c=1}^{K} p_c^2,$$

   where p_c is the proportion of samples belonging to class c in node $\mathcal{Y}$ (Breiman et al., 1984).

   - **Information Gain** evaluates the reduction in node impurity and is computed as

   $$\text{IG}(f,s) = H(\mathcal{Y}_v) - \sum_{i \in \{L,R\}} \frac{n_i}{n} H(\mathcal{Y}_i),$$

   where $H(\mathcal{Y})$ denotes the impurity of a node (entropy or Gini), n_i is the number of samples in the child node, and n is the number of samples in the parent node [Breiman et al. (1984)[.

   - **Gain Ratio** further adjusts Information Gain by penalizing splits that create many small partitions (Quinlan, 2014) and is defined as

   $$\text{GR}(f,s) = \frac{\text{IG}(f,s)}{\text{SI}(f,s)},$$

   where the split information is

   $$\text{SI}(f,s) = - \sum_{i \in \{L,R\}} \frac{n_i}{n} \log_2 \left(\frac{n_i}{n} \right).$$

3. **Evaluates Pruning Criteria**: To ensure the tree remains interpret-able and avoids over-fitting, integrated pruning is performed using cost-complexity analysis. The cost of a sub-tree T_v is given by

   $$R_\alpha(T_v) = R(T_v) + \alpha \|T_v\|,$$

   where $R(T_v)$ is the miss-classification error, $|T_v|$ is the number of leaves in the sub-tree, and $\alpha \geq 0$ is a complexity parameter controlling the trade-off between accuracy and simplicity. For a single leaf node t_v, the cost is

   $$R_\alpha(t_v) = R(t_v) + \alpha.$$

   A sub-tree T_v is pruned, i.e., replaced by a single leaf node, if

   $$R_\alpha(t_v) \leq R_\alpha(T_v),$$

   ensuring that further splits are retained only when they yield sufficient reduction in classification error to justify the added complexity.

The custom uni variate decision tree implementation in R combines axis-parallel splits with integrated cost-complexity pruning, offering several key differences from standard CART/**rpart** trees. Numeric features are split only at midpoints where class labels change, while categorical features use binary splits at each level, reducing computational complexity. Splitting criteria include Gini impurity, information gain, and Gain Ratio—the latter penalizing splits with high split information to favor balanced partitions. Integrated pruning is applied during tree construction using a fixed α parameter, replacing sub-trees with leaves whenever the cost-complexity criterion is satisfied, unlike CART/**rpart**,

---

Algorithm 1 Decision Tree with Integrated Cost-Complexity Pruning

---

1: **Procedure** GenerateTree($\mathcal{Y}, \mathbf{X}, d, d_{\max}, \alpha$)
2: $n \leftarrow |\mathcal{Y}|$
3: $\mathbf{p} \leftarrow$ ClassProbabilities($\mathcal{Y}$)
4: $\hat{y} \leftarrow \arg\max_c p_c$
5: **if** $d \geq d_{\max}$ **or** $|\text{unique}(\mathcal{Y})| \leq 1$ **then**
6:    **return** LeafNode($\hat{y}, \mathbf{p}, n$)
7: **end if**
8: $(f^*, s^*, \text{score}) \leftarrow$ FeatureSelector($\mathcal{Y}, \mathbf{X}, \text{criterion}$)
9: **if** $f^* = $ null **or** score $= -\infty$ **then**
10:    **return** LeafNode($\hat{y}, \mathbf{p}, n$)
11: **end if**
12: $\mathcal{I}_L, \mathcal{I}_R \leftarrow$ Split($\mathbf{X}[f^*], s^*$)
13: **if** $|\mathcal{I}_L| = 0$ **or** $|\mathcal{I}_R| = 0$ **then**
14:    **return** LeafNode($\hat{y}, \mathbf{p}, n$)
15: **end if**
16: $T_L \leftarrow$ GENERATETREE($\mathcal{Y}[\mathcal{I}_L], \mathbf{X}[\mathcal{I}_L], d+1, d_{\max}, \alpha$)
17: $T_R \leftarrow$ GENERATETREE($\mathcal{Y}[\mathcal{I}_R], \mathbf{X}[\mathcal{I}_R], d+1, d_{\max}, \alpha$)
18: $T_v \leftarrow$ InternalNode($f^*, s^*, T_L, T_R, \hat{y}, \mathbf{p}, n$)
19: {Integrated pruning decision}
20: $R_{\text{subtree}} \leftarrow$ SubtreeError(T_v) $+ \alpha \cdot$ CountLeaves(T_v)
21: $R_{\text{leaf}} \leftarrow (1 - \max(\mathbf{p})) + \alpha$
22: **if** $R_{\text{leaf}} \leq R_{\text{subtree}}$ **then**
23:    **return** LeafNode($\hat{y}, \mathbf{p}, n$) {Prune subtree}
24: **else**
25:    **return** T_v
26: **end if**
27: **End Procedure**

---

which performs post-pruning via cross-validation. This single-pass pruning approach reduces training time and space requirements, while still producing interpret-able trees. Predictions follow standard tree traversal, returning majority-class labels or probability estimates at leaves.

Support Vector Machine based Oblique Decision Trees

Algorithm Description

The SVM Oblique Decision Tree (SVMODT) constructs a binary classification decision tree where each internal node v at depth d. The algorithm is implemented in R using the **e1071** performs the following operations:

1. **Feature Selection:** A subset of m_d features is dynamically selected from the available feature set $\mathcal{F}$. This selection can be done using one of several strategies, including random selection, sampling, mutual information ranking (using the **FSelectorRcpp** package), or correlation-based selection. This allows the tree to adaptive-ly focus on the most informative features at each node.

2. **Feature Scaling:** Once the features are selected, z-score normalization is applied to standardize them. For the selected feature matrix $\mathbf{X}_v \in \mathbb{R}^{n \times m_d}$, the scaled features are computed as

$$\mathbf{X}_v^{\text{scaled}} = \frac{\mathbf{X}_v - \boldsymbol{\mu}_v}{\sigma_v},$$

where $\boldsymbol{\mu}_v$ and σ_v denote the mean and standard deviation of the features in node v.

3. **SVM Training:** A linear SVM is then trained on the scaled features. Optional class weights w_c can be applied to handle imbalanced data. The SVM optimization problem is formulated as

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^{n} w_{y_i} \xi_i,$$

where C is the regularization parameter, ξ_i are slack variables, and w_{y_i} represents the class-

specific weight for sample i.

4. **Node Splitting:** After training, the decision values for each sample are computed as $f(x_i) = \mathbf{w}^T x_i + b$. Samples are then partitioned into left and right child nodes according to the sign of $f(x_i)$:

   - Left child: $\{i : f(x_i) > 0\}$
   - Right child: $\{i : f(x_i) \leq 0\}$

This process recursively continues for each child node until stopping criteria are met.

Algorithm 2 SVM Oblique Decision Tree Construction

1: **Procedure** SVMSplit($\mathcal{D}, d, d_{\max}, n_{\min}$)
2: $n \leftarrow |\mathcal{D}|, \mathcal{Y} \leftarrow \text{labels}(\mathcal{D})$
3: **if** $d > d_{\max}$ **or** $|\text{unique}(\mathcal{Y})| = 1$ **or** $n < n_{\min}$ **then**
4:     **return** LeafNode($\mathcal{Y}$) {Stopping criteria}
5: **end if**
6: $m_d \leftarrow \text{DynamicMaxFeatures}(d, m_{\text{base}}, \text{strategy})$
7: $\mathcal{F}_d \leftarrow \text{SelectFeatures}(\mathcal{D}, m_d, \text{method})$
8: $\mathbf{X}_{\text{scaled}}, \text{scaler} \leftarrow \text{Scale}(\mathcal{D}[\mathcal{F}_d])$
9: **if** $|\mathcal{F}_d| = 0$ **then**
10:     **return** LeafNode($\mathcal{Y}$) {No valid features}
11: **end if**
12: $w_c \leftarrow \text{CalculateClassWeights}(\mathcal{Y}, \text{strategy})$
13: SVM $\leftarrow \text{FitLinearSVM}(\mathbf{X}_{\text{scaled}}, \mathcal{Y}, w_c)$
14: **if** SVM $= \text{null}$ **then**
15:     **return** LeafNode($\mathcal{Y}$) {SVM fitting failed}
16: **end if**
17: $\mathbf{f} \leftarrow \text{SVM.DecisionValues}(\mathbf{X}_{\text{scaled}})$
18: $\mathcal{I}_L \leftarrow \{i : f_i > 0\}, \mathcal{I}_R \leftarrow \{i : f_i \leq 0\}$
19: **if** $|\mathcal{I}_L| = 0$ **or** $|\mathcal{I}_R| = 0$ **then**
20:     **return** LeafNode($\mathcal{Y}$) {Ineffective split}
21: **end if**
22: left $\leftarrow \text{SVMSPLIT}(\mathcal{D}[\mathcal{I}_L], d + 1, d_{\max}, n_{\min})$
23: right $\leftarrow \text{SVMSPLIT}(\mathcal{D}[\mathcal{I}_R], d + 1, d_{\max}, n_{\min})$
24: **return** InternalNode(SVM, $\mathcal{F}_d$, scaler, left, right)
25: **End Procedure**

Key Hyper-parameters

1. **Feature Selection:** The algorithm supports three strategies for selecting features at each node. The maximum number of features considered at depth d is defined as:

$$m_d = \begin{cases} m_{\text{base}} & \text{constant strategy} \\ \lfloor m_{\text{base}} \cdot \alpha^{d-1} \rfloor & \text{decrease strategy} \\ \text{Uniform}(\lfloor p \cdot \ell_{\min} \rfloor, \lfloor p \cdot \ell_{\max} \rfloor) & \text{random strategy} \end{cases}$$

Here, $\alpha \in (0, 1]$ is the decrease rate for the "decrease strategy," p is the total number of features, $\ell_{\min}$ and $\ell_{\max}$ define the fractional bounds for the random strategy, and m_{base} is the base number of features at the first depth.

2. **Feature Penalty:** To encourage diversity in feature usage across the tree, previously used features have their selection weight reduced by a factor of $(1 - \lambda)$, where $\lambda \in [0, 1)$.

3. **Class Weight Handling:** To address class imbalance, the algorithm allows four weighting schemes:

   - **None:** $w_c = 1$ for all classes.
   - **Balanced:** $w_c = \frac{n}{K \cdot n_c}$, where K is the total number of classes and n_c is the number of samples in class c.
   - **Balanced sub-sample:** $w_c = \frac{1}{n_c} \cdot \frac{K}{\sum_{c'} 1/n_{c'}}$, which adjusts weights for sub-sampled data.
   - **Custom:** Users can define their own class weights.

Interaction with SVM Training and Node Splitting

The hyper-parameters described above directly influence the behavior of the SVM-based decision tree at each node. The feature selection strategy determines which subset of features $\mathcal{F}_d$ is used to fit the linear SVM at depth d, which affects the orientation and effectiveness of the decision hyper-plane. Feature penalties ensure that no single feature dominates multiple splits, promoting diversity in the learned splits and improving generalization. Once the SVM is trained on the selected and scaled features, the decision values are used to partition the samples into left and right child nodes. By controlling the number of features, penalizing repeated use, and adjusting class weights, the tree can achieve a balance between predictive accuracy, interpret-ability, and computational efficiency.

4 Data

In this study, we evaluate the SVM-based oblique decision tree (SVMODT) on two widely used benchmark datasets: **Palmer Penguins** and **Wisconsin Diagnostic Breast Cancer (WDBC)**. These datasets provide a mix of multi-class and binary classification problems, as well as numeric and categorical features, making them suitable for testing the flexibility and interpretability of SVMODTs.

Palmer Penguins

The **palmerpenguins** data set is a multi-class data set containing morphological measurements for three penguin species (*Adelie*, *Chinstrap*, and *Gentoo*) observed on the Palmer Archipelago, Antarctica. The data set includes 344 complete observations with the following features:

- **bill_length_mm**: Length of the bill (numeric)

- **bill_depth_mm**: Depth of the bill (numeric)

- **flipper_length_mm**: Length of the flipper (numeric)

- **body_mass_g**: Body mass (numeric)

- **sex**: Sex of the penguin (categorical)

- **species**: Target class (categorical with three levels)

Wisconsin Diagnostic Breast Cancer (WDBC)

The **WDBC** (Street et al., 1993) data set is a binary classification data set derived from fine needle aspirates of breast tissue. It contains 569 observations with 30 numeric features computed from digitized images of cell nuclei. The features include measures such as **radius, texture, perimeter, area, smoothness, compactness, concavity, concave points, symmetry, fractal dimension** The target variable is **diagnosis** (Malignant (M) or benign (B) tumor). This data set is widely used for testing high-dimensional binary classification algorithms. Its numeric nature allows direct evaluation of SVM splits and the effect of class weighting strategies on imbalanced classes.

Data Preparation

For both data sets, the following pre-processing steps are applied before training the SVMODT:

1. **Missing Values**: Rows with missing values are removed.
2. **Class Filtering**: For Palmer Penguins, rows corresponding to the *Gentoo* species are removed using **dplyr** to create a binary classification problem.
3. **Categorical Encoding**: Categorical variables are encoded as factor levels suitable for training.

These pre-processing steps ensure that SVMODT can learn meaningful oblique splits across both data sets while maintaining interpret-ability.

5 Usage and Examples

The svmodt package provides an intuitive interface for building SVM-based oblique decision trees. This section demonstrates the core functionality through two practical examples: ecological species classification and medical diagnosis.

Basic Usage

The primary function **svm_split()** constructs an oblique decision tree using the following key parameters:

- **data**: A data frame containing predictors and response
- response: Name of the response variable (as a string)
- **max_depth**: Maximum tree depth (controls model complexity)
- **min_samples**: Minimum samples required to split a node
- **max_features**: Maximum features to consider at each split
- **feature_method**: Feature selection strategy ("random", "mutual", "cor")
- **class_weights**: Strategy for handling class imbalance

Example 1: Penguin Species Classification

We demonstrate the package using the **palmerpenguins** data set, which contains morphological measurements for three penguin species. We focus on classifying Adelie and Chinstrap penguins.

Model Training

In this example, the model is trained on the train_data data set to predict the categorical response variable, *species*. The function svm_split() recursively partitions the data using **linear SVM hyperplanes** at each internal node rather than conventional univariate thresholds, thereby producing **oblique splits** capable of capturing multivariate relationships among predictors. The argument max_depth = 3 constrains the tree to a maximum of three levels, providing control over model complexity and mitigating overfitting. Similarly, max_features = 2 restricts each node to consider only two features when fitting the SVM, enhancing computational efficiency and interpretability. The parameter feature_method = "mutual" specifies that feature selection at each node is based on **mutual information**, ensuring that the most informative variables relative to the response are prioritized. Setting verbose = FALSE suppresses intermediate output for streamlined execution.

```
# Train SVMODT with mutual information feature selection
tree <- svm_split(
  data = train_data,
  response = "species",
  max_depth = 3,
  max_features = 2,
  feature_method = "mutual",
  verbose = FALSE
)
```

Model Structure

Each internal node represents a **binary linear decision boundary** obtained from a fitted SVM model using the features listed at that node. For example, the root node (depth = 1) uses bill_length_mm and flipper_length_mm to form the first separating hyper-plane. Observations satisfying $SVM > 0$ proceed to the left child, while those with $SVM \leq 0$ move to the right branch.

```
#> [Node] depth = 1 | n = 175 | features = [bill_length_mm,flipper_length_mm] | max_feat = 2
#>  |- Left branch (SVM > 0):
#> |   [Node] depth = 2 | n = 125 | features = [bill_length_mm,bill_depth_mm]
#> |   |- Left branch (SVM > 0):
#> |   |   [Leaf] predict = Adelie | n = 123 | features = [bill_length_mm,bill_depth_mm]
```

Figure 1: Visualization of the root node (depth = 1) of the SVM-based oblique decision tree.

```
#> |    `- Right branch (SVM <= 0):
#> |       (no right child)
#>  `- Right branch (SVM <= 0):
#>    [Node] depth = 2 | n = 50 | features = [bill_length_mm,bill_depth_mm]
#>    |- Left branch (SVM > 0):
#>    |   [Leaf] predict = Chinstrap | n = 48
#>    `- Right branch (SVM <= 0):
#>       (no right child)
```

Terminal nodes (labeled as *Leaf*) correspond to final class predictions. In this case, most observations in the left sub-tree are classified as *Adelie*, while the right sub-tree primarily identifies *Chinstrap* penguins. The structure illustrates how **SVM-based oblique splits** allow the tree to capture complex linear interactions between multiple predictors—something that traditional axis-aligned decision trees cannot achieve.

Visualizing Decision Boundaries

The package includes visualization tools to examine the SVM hyper-planes at each node, allowing users to interpret how features contribute to class separation within the tree. *Note that visualization is currently supported only for trees where each node consistently uses exactly two features,* ensuring that the separating hyper-plane can be rendered clearly in two-dimensional feature space. The `visualize_svm_tree()` function produces a graphical representation of the overall decision hierarchy

Fig 1 depicts the root node of the SVM-based oblique decision tree, where the first split is performed using the features `bill_length_mm` and `flipper_length_mm`. The figure illustrates how the linear SVM at the root node divides the dataset into two branches, guiding samples toward subsequent child nodes. **Note:** This visualization is only possible because the node considers exactly two features, a requirement for 2D plotting.

Similarly, Fig 2 shows a node at depth 2 of the tree, where the SVM uses bill_length_mm and bill_depth_mm to further partition the data. The decision hyper-plane and sample positions are plotted, highlighting how oblique splits can capture multivariate relationships that univariate thresholds cannot. This figure demonstrates the recursive nature of the tree and the refinement of class separation
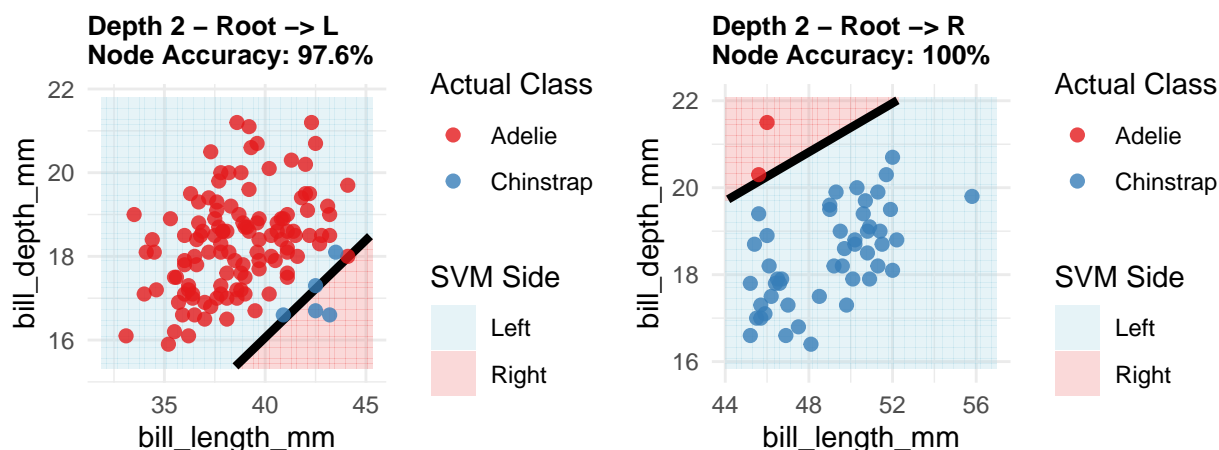
Figure 2: Visualization of a node (depth = 2).

at deeper levels. As with the root node, the visualization is restricted to nodes consistently using exactly two features.

Prediction Path Tracing

The svmodt package allows tracing the prediction path of individual observations through the SVM-based oblique decision tree. For each sample, the relevant feature values are evaluated at each node using the linear SVM hyper-plane associated with that node. For instance, a sample with `bill_length_mm = 39.5`, `bill_depth_mm = 17.4`, `flipper_length_mm = 186`, and `body_mass_g = 3800` first reaches the root node, which considers `bill_length_mm` and `flipper_length_mm`. The SVM decision value at this node is `2.1098`, which directs the sample along the left branch. At depth 2, the node evaluates `bill_length_mm` and `bill_depth_mm` and produces a decision value of `1.6592`, again guiding the sample to the left branch. Finally, the sample reaches a leaf node that predicts the class `Adelie`, which contains 123 training samples. The complete traversal path is LEFT -> LEFT, resulting in the final prediction of `Adelie`. This tracing functionality provides interpret-able insights into how each decision is made within the oblique decision tree.

```
#> === Tracing Prediction Path ===
#> Sample 1 :
#>    species = 1
#>    bill_length_mm = 39.5
#>    bill_depth_mm = 17.4
#>    flipper_length_mm = 186
#>    body_mass_g = 3800
#>
#>  [Node 1 ] features = bill_length_mm,flipper_length_mm
#>    SVM decision value: 2.1098
#>    -> Going LEFT (decision > 0)
#>    [Node 2 ] features = bill_length_mm,bill_depth_mm
#>     SVM decision value: 1.6592
#>     -> Going LEFT (decision > 0)
#>     [FINAL] Predict Adelie (n = 123 )
#>     Path taken: LEFT -> LEFT
#>
#> Final prediction: Adelie

#> [1] "Adelie"
```

Example 2: Breast Cancer Diagnosis

We now demonstrate the package on the Wisconsin Diagnostic Breast Cancer (WDBC) data set, highlighting advanced functionalities such as handling class imbalance, feature sub-setting, and feature penalization. These features allow the SVM-based oblique decision tree to effectively manage real-world challenges, including unequal class distributions, high-dimensional feature spaces, and repeated feature usage, while maintaining interpret-ability and predictive performance.

Handling Class Imbalance

The package provides built-in support for handling class imbalance through configurable class weighting. In the first example, the svm_split() function is applied to the WDBC data set to predict the binary response variable with class_weights = "balanced", which automatically adjusts the weight of each class inversely proportional to its frequency. This ensures that minority classes receive higher influence during SVM training, improving predictive performance on imbalanced data sets.

```
# Train with balanced class weights
tree_balanced <- svm_split(
  data = train_wdbc,
  response = "diagnosis",
  max_depth = 4,
  max_features = 5,
  feature_method = "mutual",
  class_weights = "balanced",  # Automatic balancing
  verbose = FALSE
)

print_svm_tree(tree = tree_balanced,
               show_feature_info = FALSE,
               show_penalties = FALSE, show_probabilities = TRUE)

#>  [Node] depth = 1 | n = 455
#>  |- Left branch (SVM > 0):
#>  |    [Node] depth = 2 | n = 177
#>  |    |- Left branch (SVM > 0):
#>  |    |   [Node] depth = 3 | n = 147
#>  |    |   |- Left branch (SVM > 0):
#>  |    |   |   [Leaf] predict = M | n = 144 | probs = [B = 0, M = 1]
#>  |    |   `- Right branch (SVM <= 0):
#>  |    |       (no right child)
#>  |    `- Right branch (SVM <= 0):
#>  |       [Node] depth = 3 | n = 30
#>  |       |- Left branch (SVM > 0):
#>  |       |   [Node] depth = 4 | n = 13
#>  |       |   |- Left branch (SVM > 0):
#>  |       |   |   [Leaf] predict = M | n = 5 | probs = [B = 0, M = 1]
#>  |       |   `- Right branch (SVM <= 0):
#>  |       |       [Leaf] predict = B | n = 8 | probs = [B = 0.5, M = 0.5]
#>  |       `- Right branch (SVM <= 0):
#>  |          [Node] depth = 4 | n = 17
#>  |          |- Left branch (SVM > 0):
#>  |          |   [Leaf] predict = B | n = 13 | probs = [B = 1, M = 0]
#>  |          `- Right branch (SVM <= 0):
#>  |              (no right child)
#>  `- Right branch (SVM <= 0):
#>     [Node] depth = 2 | n = 278
#>     |- Left branch (SVM > 0):
#>     |   [Leaf] predict = B | n = 256 | probs = [B = 0.988, M = 0.012]
#>     `- Right branch (SVM <= 0):
#>        [Node] depth = 3 | n = 22
#>        |- Left branch (SVM > 0):
#>        |   [Leaf] predict = M | n = 5 | probs = [B = 0.4, M = 0.6]
#>        `- Right branch (SVM <= 0):
#>           [Node] depth = 4 | n = 17
#>           |- Left branch (SVM > 0):
#>           |   [Leaf] predict = B | n = 11 | probs = [B = 1, M = 0]
#>           `- Right branch (SVM <= 0):
#>               [Leaf] predict = B | n = 6 | probs = [B = 0.833, M = 0.167]
```

In the second example, custom class weights can be specified via the custom_class_weights argument, assigning a higher penalty to the malignant class ("M" = 3) relative to the benign class ("B" = 1) to emphasize minimizing false negatives. The resulting trees (tree_balanced and tree_custom)

incorporate these weights during the recursive SVM-based splits, enabling more equitable class separation while still leveraging the oblique decision tree structure.

```
# Custom class weights for domain-specific costs
custom_weights <- c("B" = 1, "M" = 3)  # Penalize false negatives
tree_custom <- svm_split(
  data = train_wdbc,
  response = "diagnosis",
  max_depth = 4,
  max_features = 5,
  class_weights = "custom",
  custom_class_weights = custom_weights,
  verbose = FALSE
)

print_svm_tree(tree = tree_custom,
               show_feature_info = FALSE,
               show_penalties = FALSE,
               show_probabilities = FALSE)

#> [Node] depth = 1 | n = 455
#> |- Left branch (SVM > 0):
#> |   [Node] depth = 2 | n = 190
#> |   |- Left branch (SVM > 0):
#> |   |   [Node] depth = 3 | n = 188
#> |   |   |- Left branch (SVM > 0):
#> |   |   |   [Leaf] predict = M | n = 171
#> |   |   `- Right branch (SVM <= 0):
#> |   |       [Leaf] predict = B | n = 17
#> |   `- Right branch (SVM <= 0):
#> |       (no right child)
#> `- Right branch (SVM <= 0):
#>     [Node] depth = 2 | n = 265
#>     |- Left branch (SVM > 0):
#>     |   [Node] depth = 3 | n = 264
#>     |   |- Left branch (SVM > 0):
#>     |   |   [Leaf] predict = B | n = 263
#>     |   `- Right branch (SVM <= 0):
#>     |       (no right child)
#>     `- Right branch (SVM <= 0):
#>         (no right child)
```

Feature Penalization

To encourage feature diversity across the tree, the package implements a feature penalization mechanism that reduces the likelihood of reusing features selected at ancestor nodes.

```
set.seed(123)
tree_penalty <- svm_split(
  data = train_wdbc,
  response = "diagnosis",
  max_depth = 4,
  max_features = 4,
  penalize_used_features = TRUE,
  feature_penalty_weight = 0.6,
  verbose = FALSE
)

stringr::str_wrap(print_svm_tree(tree_penalty, show_probabilities = FALSE,
               show_feature_info = TRUE,
               show_penalties = FALSE))

#> [Node] depth = 1 | n = 455 | features = [smoothness_se,symmetry_se,area_se,perimeter_mean] | max_feat = 4
```

```
#>  |- Left branch (SVM > 0):
#> |  [Leaf] predict = M | n = 138 | features = [perimeter_mean,perimeter_se,fractal_dimension_worst,compactness_
#>  `- Right branch (SVM <= 0):
#>   [Node] depth = 2 | n = 317 | features = [concave.points_mean,radius_se,compactness_mean,area_worst] | max_fe
#>     |- Left branch (SVM > 0):
#>   |  [Leaf] predict = B | n = 291 | features = [perimeter_worst,concave.points_se,concavity_se,texture_mean]
#>     `- Right branch (SVM <= 0):
#>     [Node] depth = 3 | n = 26 | features = [concavity_worst,perimeter_mean,smoothness_mean,compactness_worst]
#>        |- Left branch (SVM > 0):
#>      |  [Leaf] predict = M | n = 22
#>        `- Right branch (SVM <= 0):
#>         (no right child)

#> character(0)
```

Penalization is controlled via the `feature_penalty_weight` parameter, which scales down the selection probability of previously used features. For instance, setting `feature_penalty_weight = 0.6` reduces the weight of used features by 60%, promoting the consideration of alternative predictors at subsequent splits. This mechanism helps improve model robustness, reduces redundancy in feature usage, and enhances interpret-ability by encouraging the tree to explore diverse feature combinations.

Dynamic Feature Selection

The package also supports dynamic feature selection strategies, which vary the number of features considered at each node depending on tree depth or randomness.

```
tree_decrease <- svm_split(
  data = train_wdbc,
  response = "diagnosis",
  feature_method = "mutual",
  max_depth = 3,
  max_features = 4,
  max_features_strategy = "decrease",
  max_features_decrease_rate = 0.5,
  verbose = FALSE
)

stringr::str_wrap(print_svm_tree(tree_decrease, show_penalties = FALSE,
                show_feature_info = TRUE))

#> [Node] depth = 1 | n = 455 | features = [perimeter_worst,concave.points_mean,concave.points_worst,radius_wors
#>  |- Left branch (SVM > 0):
#> |   [Node] depth = 2 | n = 155 | features = [perimeter_worst,texture_worst]
#> |   |- Left branch (SVM > 0):
#> |  |   [Leaf] predict = M | n = 152 | features = [concavity_worst]
#> |   `- Right branch (SVM <= 0):
#> |       (no right child)
#>  `- Right branch (SVM <= 0):
#>     [Leaf] predict = B | n = 300 | features = [radius_worst,perimeter_worst]

#> character(0)
```

In the **decrease strategy**, the number of candidate features at depth d is scaled by a decay factor α, e.g., $m_0 = m_0 \cdot \alpha^{d-1}$, reducing feature count at deeper levels to improve efficiency and prevent over-fitting.

```
set.seed(123)

tree_random <- svm_split(
  data = train_wdbc,
  response = "diagnosis",
  feature_method = "mutual",
  max_depth = 4,
```

```
  max_features_strategy = "random",
  max_features_random_range = c(0.05, 0.2),  # 30-80% of features
  verbose = FALSE
)

print_svm_tree(tree_random, show_penalties = FALSE,
               show_feature_info = TRUE)

#> [Node] depth = 1 | n = 455 | features = [perimeter_worst,concave.points_mean,concave.points_worst,radius_wors
#>  |- Left branch (SVM > 0):
#> |   [Node] depth = 2 | n = 155 | features = [perimeter_worst,texture_worst]
#> |   |- Left branch (SVM > 0):
#> | |   [Node] depth = 3 | n = 152 | features = [concavity_worst,concavity_mean,concave.points_mean,concavity_se
#> | |   |- Left branch (SVM > 0):
#> | |   |   [Leaf] predict = M | n = 149 | features = [radius_mean,texture_mean]
#> | |   `- Right branch (SVM <= 0):
#> | |       (no right child)
#> |   `- Right branch (SVM <= 0):
#> |       (no right child)
#>  `- Right branch (SVM <= 0):
#>    [Node] depth = 2 | n = 300 | features = [radius_worst,perimeter_worst,area_worst,area_mean,area_se]
#>     |- Left branch (SVM > 0):
#>     | [Leaf] predict = B | n = 299 | features = [area_mean,radius_worst,perimeter_worst,area_worst]
#>     `- Right branch (SVM <= 0):
#>         (no right child)
```

Alternatively, the **random strategy** selects a feature count uniformly at random from a specified range (e.g., 5–20% of the total features) at each node, promoting stochasticity and diversity in splits. These mechanisms provide flexible control over tree complexity and feature exploration, particularly useful in high-dimensional data sets such as the WDBC data set.

6 Results

7 Discussion

Bibliography

M. Bala and R. Agrawal. Optimal decision tree based multi-class support vector machine. *Informatica*, 35(2), 2011. [p4]

K. Bennett and J. Blue. A support vector machine approach to decision trees. In *1998 IEEE International Joint Conference on Neural Networks Proceedings. IEEE World Congress on Computational Intelligence (Cat. No.98CH36227)*, volume 3, pages 2396–2401 vol.3, 1998. doi: 10.1109/IJCNN.1998.687237. [p4]

L. Breiman, J. H. Friedman, R. A. Olshen, and C. J. Stone. *Classification and Regression Trees*. Wadsworth, 1984. URL http://lyle.smu.edu/~mhd/8331f06/cart.pdf. [p1, 5]

C. E. Brodley and P. E. Utgoff. Multivariate decision trees. *Machine Learning*, 19:45–77, 1995. URL https://api.semanticscholar.org/CorpusID:16836572. [p1, 2]

L. Cañete, R. Monroy, and M. Medina-Pérez. A review and experimental comparison of multivariate decision trees. *IEEE Access*, PP:1–1, 08 2021. doi: 10.1109/ACCESS.2021.3102239. [p2]

J. Cervantes, F. Garcia-Lamont, L. Rodríguez-Mazahua, and A. Lopez. A comprehensive survey on support vector machine classification: Applications, challenges and trends. *Neurocomputing*, 408:189–215, 2020. ISSN 0925-2312. doi: https://doi.org/10.1016/j.neucom.2019.10.118. URL https://www.sciencedirect.com/science/article/pii/S0925231220307153. [p3, 4]

C. Cortes and V. Vapnik. Support-vector networks. *Machine learning*, 20(3):273–297, 1995. [p2]

N. Cristianini. An introduction to support vector machines and other kernel-based learning methods, 2000. [p2]

J.-x. Dong, A. Krzyzak, and C. Y. Suen. Fast svm training algorithm with decomposition on very large data sets. *IEEE transactions on pattern analysis and machine intelligence*, 27(4):603–618, 2005. [p3]

M. R. Frean. *Small nets and short paths: Optimising neural computation*. PhD thesis, 1990. [p2]

M. Friedl and C. Brodley. Decision tree classification of land cover from remotely sensed data. *Remote Sensing of Environment*, 61(3):399–409, 1997. ISSN 0034-4257. doi: https://doi.org/10.1016/S0034-4257(97)00049-7. URL https://www.sciencedirect.com/science/article/pii/S0034425797000497. [p2]

S. I. Gallant. Optimal linear discriminants. *Eighth International Conference on Pattern Recognition*, pages 849–852, 1986. URL https://cir.nii.ac.jp/crid/1573668924476144256. [p2]

M. Ganaie, M. Tanveer, P. Suganthan, and V. Snasel. Oblique and rotation double random forest. *Neural Networks*, 153:496–517, 2022. ISSN 0893-6080. doi: https://doi.org/10.1016/j.neunet.2022.06.012. URL https://www.sciencedirect.com/science/article/pii/S0893608022002258. [p4]

C.-W. Hsu and C.-J. Lin. A comparison of methods for multiclass support vector machines. *IEEE transactions on Neural Networks*, 13(2):415–425, 2002. [p4]

J. Kittler. Feature selection and extraction. *Handbook of Pattern Recognition and Image Processing*, pages 59–83, 1986. URL https://cir.nii.ac.jp/crid/1573950400671608448. [p2]

M. Koziol and M. Wozniak. *Multivariate Decision Trees vs. Univariate Ones*, pages 275–284. Springer Berlin Heidelberg, Berlin, Heidelberg, 2009. ISBN 978-3-540-93905-4. doi: 10.1007/978-3-540-93905-4_33. URL https://doi.org/10.1007/978-3-540-93905-4_33. [p2]

H. Liu and R. Setiono. Feature transformation and multivariate decision tree induction. In S. Arikawa and H. Motoda, editors, *Discovey Science*, pages 279–291, Berlin, Heidelberg, 1998. Springer Berlin Heidelberg. ISBN 978-3-540-49292-4. [p2]

O. Mangasarian and E. Wild. Multisurface proximal support vector machine classification via generalized eigenvalues. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 28(1):69–74, 2006. doi: 10.1109/TPAMI.2006.17. [p4]

V. Menkovski, I. T. Christou, and S. Efremidis. Oblique decision trees using embedded support vector machines in classifier ensembles. In *2008 7th IEEE International Conference on Cybernetic Intelligent Systems*, pages 1–6, 2008. doi: 10.1109/UKRICIS.2008.4798937. [p4]

J. Mingers. An empirical comparison of pruning methods for decision tree induction. *Machine Learning*, 4:227–243, 01 1989. doi: 10.1023/A:1022604100933. [p1]

R. Montañana, J. A. Gámez, and J. M. Puerta. Stree: a single multi-class oblique decision tree based on support vector machines, 2021. [p4]

R. Montañana, J. A. Gámez, and J. M. Puerta. Odte—an ensemble of multi-class svm-based oblique decision trees. *Expert Systems with Applications*, 273:126833, 2025. ISSN 0957-4174. doi: https://doi.org/10.1016/j.eswa.2025.126833. URL https://www.sciencedirect.com/science/article/pii/S0957417425004555. [p4]

M. Nanda, K. Seminar, D. Nandika, and A. Maddu. A comparison study of kernel functions in the support vector machine and a comparison study of kernel functions in the support vector machine and its application for termite detection, 2018. [p3]

G. Pagallo and D. Haussler. Boolean feature discovery in empirical learning. *Machine Learning*, 5:71–99, 1990. URL https://api.semanticscholar.org/CorpusID:5661437. [p1]

F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011. [p4]

J. R. Quinlan. Induction of decision trees. *Machine Learning*, 1:81–106, 1986. URL https://api.semanticscholar.org/CorpusID:189902138. [p1]

J. R. Quinlan. *C4. 5: programs for machine learning*. Elsevier, 2014. [p1, 5]

R. Rivera-Lopez and J. Canul-Reich. Construction of near-optimal axis-parallel decision trees using a differential-evolution-based approach. *IEEE Access*, PP:1–1, 01 2018. doi: 10.1109/ACCESS.2017.2788700. [p1]

C. Schaffer. Deconstructing the digit recognition problem. In D. Sleeman and P. Edwards, editors, *Machine Learning Proceedings 1992*, pages 394–399. Morgan Kaufmann, San Francisco (CA), 1992. ISBN 978-1-55860-247-2. doi: https://doi.org/10.1016/B978-1-55860-247-2.50056-5. URL https://www.sciencedirect.com/science/article/pii/B9781558602472500565. [p1]

W. N. Street, W. H. Wolberg, and O. L. Mangasarian. Nuclear feature extraction for breast tumor diagnosis. In *Electronic imaging*, 1993. URL https://api.semanticscholar.org/CorpusID:14922543. [p8]

F. TAKAHASHI and S. Abe. Decision-tree-based multiclass support vector machines. volume 3, pages 1418 – 1422 vol.3, 12 2002. ISBN 981-04-7524-1. doi: 10.1109/ICONIP.2002.1202854. [p4]

V. Vapnik. *The nature of statistical learning theory*. Springer science & business media, 2013. [p3, 4]

P. C. Young. Recursive estimation and time-series analysis: An introduction. 1984. URL https://api.semanticscholar.org/CorpusID:60181335. [p2]

Aneesh Agarwal
Monash University

aaga0022@student.monash.edu

Jack Jewson
Monash University
Department of Econometrics and Business Statistics, Monash University, Australia
Jack.Jewson@monash.edu

Erik Sverdrup
Monash University
Department of Econometrics and Business Statistics, Monash University, Australia
Erik.Sverdrup@monash.edu